

CS57300
PURDUE UNIVERSITY
FEBRUARY 19, 2019

DATA MINING

ANNOUNCEMENT

- ▶ Assignment 3 is out
 - ▶ Implement Logistic Regression and Linear SVM for speed dating event outcome prediction
 - ▶ Due: March 8 (Friday), 11:59pm
- ▶ In-class midterm exam in two weeks
 - ▶ March 5, 4:30-5:45pm, WANG 2599

SMOOTH OPTIMIZATION

SOLVE CONVEX OPTIMIZATION PROBLEM

- ▶ Minimize a convex function without any constraints on the variables
 - ▶ If $f'(x)=0$ then x is a stationary point of f
 - ▶ If $f'(x)=0$ and $f''(x)$ is not negative then x is a local minimum of f (for convex function, this is also a global minimum)
 - ▶ If f is a strictly convex function, any stationary point of f is the unique global minimum of f
- ▶ What about minimizing a convex function with constraints?

USE LAGRANGE MULTIPLIERS TO SOLVE CONVEX OPTIMIZATION

- ▶ For a standard form of convex optimization problem (f_0 and f_i are convex, h_i is linear):

$$\begin{array}{ll}\text{minimize} & f_0(x) \\ \text{subject to} & f_i(x) \leq 0, \quad \text{for } i = 1, \dots, m. \\ & h_i(x) = 0, \quad \text{for } i = 1, \dots, k.\end{array}$$

- ▶ The Lagrangian function of it is

$$L(x, \lambda, v) = f_0(x) + \sum_{i=1}^m \lambda_i f_i(x) + \sum_{i=1}^k v_i h_i(x)$$

- ▶ $\lambda_i \geq 0$ is the **Lagrange multiplier** for the i -th inequality constraint, v_i is the **Lagrange multiplier** for the i -th equality constraint
- ▶ Solve the constrained optimization problem by finding the stationary point of the Lagrangian function

LOGISTIC REGRESSION LEARNING

- ▶ Logistic regression: $P(y = 1|\mathbf{x}) = \frac{1}{1 + e^{-(\mathbf{w}^T \mathbf{x} + w_0)}}$
- ▶ Maximize (log) likelihood: $\mathbf{w} = (\mathbf{w}, w_0), \mathbf{x}_i = (\mathbf{x}_i, 1)$

$$\begin{aligned} \log L(\mathbf{w} | D) &= \sum_{i=1}^N \log p(y_i | \mathbf{w}) \\ &= \sum_{i=1}^N \log \left[\left(\frac{1}{1 + e^{-\mathbf{w}^T \mathbf{x}_i}} \right)^{y_i} \left(\frac{e^{-\mathbf{w}^T \mathbf{x}_i}}{1 + e^{-\mathbf{w}^T \mathbf{x}_i}} \right)^{1-y_i} \right] \\ &= \sum_{i=1}^N (y_i \mathbf{w}^T \mathbf{x}_i - \log(1 + e^{\mathbf{w}^T \mathbf{x}_i})) \end{aligned}$$

- ▶ Minimize: $\sum_{i=1}^N (-y_i \mathbf{w}^T \mathbf{x}_i + \log(1 + e^{\mathbf{w}^T \mathbf{x}_i}))$

LOGISTIC REGRESSION LEARNING

$$\text{minimize } \sum_{i=1}^N (-y_i \mathbf{w}^T \mathbf{x}_i + \log(1 + e^{\mathbf{w}^T \mathbf{x}_i}))$$

$$\begin{aligned} \frac{d \log L(\mathbf{w} | D)}{d w_j} &= \sum_{i=1}^N \left(-y_i x_{ij} + \frac{1}{1 + e^{\mathbf{w}^T \mathbf{x}_i}} e^{\mathbf{w}^T \mathbf{x}_i} x_{ij} \right) \\ &= \sum_{i=1}^N \left(-y_i + \frac{1}{1 + e^{\mathbf{w}^T \mathbf{x}_i}} e^{\mathbf{w}^T \mathbf{x}_i} \right) x_{ij} \\ &= \sum_{i=1}^N (-y_i + P(y_i = 1 | \mathbf{w})) x_{ij} \end{aligned}$$

Convex!

But no closed form solution!

GRADIENT DESCENT

- ▶ For some convex functions, we may be able to take the derivative, but it may be difficult to directly solve for parameter values
- ▶ Solution:
 - ▶ Start at some value of the parameters
 - ▶ Take derivative and use it to move the parameters in the direction of the negative gradient
 - ▶ Repeat until stopping criteria is met (e.g., gradient close to 0)

Gradient Descent Rule:

$$\underline{\mathbf{w}}_{\text{new}} = \underline{\mathbf{w}}_{\text{old}} - \eta \Delta(\underline{\mathbf{w}})$$

where

$\Delta(\underline{\mathbf{w}})$ is the gradient and
 η is the learning rate (small, positive)

Notes:

1. This moves us downhill in direction $\Delta(\underline{\mathbf{w}})$ (steepest downhill direction)
2. How far we go is determined by the value of η

ILLUSTRATION OF GRADIENT DESCENT




ILLUSTRATION OF GRADIENT DESCENT






ILLUSTRATION OF GRADIENT DESCENT



**Direction of steepest
descent = direction of
negative gradient**

ILLUSTRATION OF GRADIENT DESCENT



For convex functions, when gradient descent converges, the solution is global minimum.

STOPPING CRITERIA FOR GRADIENT DESCENT


- ▶ Ideally, $f'(x)=0\dots$
- ▶ In practice...
 - ▶ $\|\nabla f(x)\| < \varepsilon$
 - ▶ $|f(x_{k+1}) - f(x_k)| < \varepsilon$
 - ▶ $\|x_{k+1} - x_k\| < \varepsilon$
 - ▶ Maximum number of iterations has been reached

GRADIENT ASCENT

- ▶ For concave functions that you want to *maximize*, take a step in direction of gradient (i.e., $w_{\text{new}} \leftarrow w_{\text{old}} + \eta \nabla(w)$)
- ▶ Otherwise same as gradient descent:
 - ▶ Start at some parameter values
 - ▶ Take derivative, move the parameters in the direction of gradient
 - ▶ Repeat until stopping criteria is met (e.g., gradient close to 0)

GRADIENT DESCENT FOR NON-CONVEX OPTIMIZATION

- ▶ Works on any objective function $F(\theta)$
 - ▶ as long as we can evaluate the gradient $\Delta(\theta)$
 - ▶ this can be very useful for minimizing complex functions F
- ▶ Can be used in hill-climbing search to find local minima in smooth, but non-convex functions



- ▶ If function has multiple local minima, gradient descent goes to the closest local minimum:
 - ▶ solution: random restarts from multiple places in model space

LOGISTIC REGRESSION: RECAP

LOGISTIC REGRESSION

- ▶ Same parametric form as standard regression, but uses logistic function for binary classification


Logistic regression model:

$$P(y = 1|\mathbf{x}) = \frac{1}{1 + e^{-(\mathbf{w}^T \mathbf{x} + w_0)}}$$

- ▶ Output is the (positive) class probability rather than the binary prediction
- ▶ Logistic function transform ensures output is $[0,1]$

Logistic function:

$$\text{logistic}(x) := \frac{e^x}{1 + e^x} = \frac{1}{1 + e^{-x}}$$



LR EXAMPLE

$$P(BC = 1|A, I, S, CR) = \frac{1}{1 + e^{-(\mathbf{w}^T \mathbf{x})}}$$

$$\mathbf{x} = [Int, A, I, S, CR]$$

$$\mathbf{w} = [w_0, w_A, w_I, w_S, w_{CR}]$$

LR parameters = w

Intercept	Age>40	Income=high	Student=yes	Credit=fair	BuysComp?
1	0	1	0	1	0
1	0	1	0	0	0
1	0	1	0	1	1
1	1	0	0	1	1
1	1	0	1	1	1
1	1	0	1	0	0
1	0	0	1	0	1
1	0	0	0	1	0
1	0	0	1	1	1
1	1	0	1	1	1
1	0	0	1	0	1
1	0	0	0	0	1
1	0	1	1	1	1
1	1	0	0	0	0

- ▶ Score function: likelihood
- ▶ Estimate \mathbf{w} with maximum likelihood estimation

LR LEARNING

- ▶ Score function: likelihood function

$$\text{minimize} \sum_{i=1}^N (-y_i \mathbf{w}^T \mathbf{x}_i + \log(1 + e^{\mathbf{w}^T \mathbf{x}_i}))$$

- ▶ Estimate optimal \mathbf{w} using gradient descent $\frac{d \log L}{d w_j} = \sum_{i=1}^N (-y_i + P(y_i = 1 | \mathbf{w}, \mathbf{x}_i)) x_{ij}$

Gradient descent:

Start at some \mathbf{w} , e.g., $\mathbf{w}=[0,0,0,0,0]$

Make predictions given current \mathbf{w} :

Calculate gradient for each parameter:

Move parameters in direction of negative gradient:

Repeat until stopping criteria is met

$$\forall i \quad \hat{y}_i = P(y_i = 1 | \mathbf{x}_i) = \frac{1}{1 + e^{-\mathbf{w}^T \mathbf{x}_i}}$$

$$\forall j \quad \frac{d \log L}{d w_j} = \left[\sum_{i=1}^n (-y_i + \hat{y}_i) x_{ij} \right] = \nabla_j$$

$$\forall j \quad w_j^{new} = w_j - \eta \nabla_j$$

LR PREDICTION

Intercept	Age>40	Income=high	Student=yes	Credit=fair	BuysComp?
1	0	1	0	1	0
1	0	1	0	0	0
1	0	1	0	1	1
1	1	0	0	1	1
1	1	0	1	1	1
1	1	0	1	0	0
1	0	0	1	0	1
1	0	0	0	1	0
1	0	0	1	1	1
1	1	0	1	1	1
1	0	0	1	0	1
1	0	0	0	0	1
1	0	1	1	1	1
1	1	0	0	0	0
1	0	1	0	0	?

- What is the probability that new person will buy a computer?

$$\mathbf{x} = [1, 0, 1, 0, 0]$$

$$\mathbf{w} = [-1.3, 1, 2, -2, 0.7]$$

$$\mathbf{x}^T \mathbf{w} = 0.7$$

$$P(BC = 1 | \mathbf{x}) = \frac{1}{1 + e^{-0.7}} = 0.668$$

DEAL WITH OVERFITTING

- ▶ Simply finding the parameter values that lead to maximum likelihood function value in the training dataset may imply overfitting!
- ▶ Solution: add a **regularization term** in the scoring function to penalize complex models
 - ▶ e.g., L2 regularization term: $\frac{\lambda}{2} \|w\|^2$
 - ▶ λ is the regularization parameter; the larger the value, the more we are in favor of simple models

LR LEARNING WITH REGULARIZATION TERM

- ▶ Score function: likelihood with L2 regularization

$$\text{minimize} \sum_{i=1}^N (-y_i \mathbf{w}^T \mathbf{x}_i + \log(1 + e^{\mathbf{w}^T \mathbf{x}_i})) + \frac{\lambda}{2} \|\mathbf{w}\|^2$$

- ▶ Estimate optimal \mathbf{w} using gradient descent

Gradient descent:

Start at some \mathbf{w} , e.g., $\mathbf{w}=[0,0,0,0,0]$

Make predictions given current \mathbf{w} :

Calculate gradient for each parameter:

$$\begin{aligned} \forall i \quad \hat{y}_i &= P(y_i = 1 | \mathbf{x}_i) = \frac{1}{1 + e^{-\mathbf{w}^T \mathbf{x}_i}} \\ \forall j \quad \frac{d \log L}{d w_j} &= \left[\sum_{i=1}^n (-y_i + \hat{y}_i) x_{ij} \right] + \lambda w_j \\ &= \nabla_j \end{aligned}$$

Move parameters in direction of negative gradient:


$$\forall j \quad w_j^{new} = w_j - \eta \nabla_j$$

Repeat until stopping criteria is met

SVM: RECAP

SVM: KNOWLEDGE REPRESENTATION AND SCORING FUNCTION

- ▶ Linear SVM: $y = \text{sign} \left[\sum_{i=1}^m w_i x_i + b \right]$
- ▶ Margin = $d_+ + d_- = 2/\|\mathbf{w}\|$
- ▶ Optimization problem
 - ▶ $\max 2/\|\mathbf{w}\|$
 - ▶ subject to
 $y_i(\mathbf{w}^T \mathbf{x}_i + b) \geq 1, \forall i \in \{1, 2, \dots, N\}$



SVM LEARNING

- ▶ Equivalent to minimize $\|\mathbf{w}\|^2/2$ subject to


$$y_i(\mathbf{w}^T \mathbf{x}_i + b) \geq 1, \forall i \in \{1, 2, \dots, N\}$$

- ▶ This is a **quadratic optimization** problem subject to linear constraints, there is a unique minimum

- ▶ Lagrangian function $L(\mathbf{w}, b, \lambda_i) = \frac{1}{2}\|\mathbf{w}\|^2 + \sum_{i=1}^N \lambda_i(1 - y_i(\mathbf{w}^T \mathbf{x}_i + b))$

WHAT ABOUT LINEARLY NON-SEPARABLE DATA?

- ▶ Introduce slack variables $\varepsilon_i \geq 0$ such that:
$$y_i(\mathbf{w}^T \mathbf{x}_i + b) \geq 1 - \varepsilon_i, \forall i \in \{1, 2, \dots, N\}$$
- ▶ ε_i measures the amount of error
 - ▶ When $0 < \varepsilon_i \leq 1$, data is between the margin, but classified correctly
 - ▶ When $\varepsilon_i > 1$, data is misclassified



“SOFT” MARGIN OPTIMIZATION

- ▶ With slack variables the score function is:

$$\min_{\mathbf{w}, \xi} \|\mathbf{w}\|^2 + C \sum_i^N \xi_i$$

- ▶ And new constraints:

$$y_i(x_i \cdot w + b) - (1 - \xi_i) \geq 0 \quad \forall i$$

- ▶ If ξ are sufficiently large, then every constraint can be satisfied
- ▶ C is regularization parameter
 - ▶ Small C means constraints can be ignored in order to find large margin
 - ▶ Large C means constraints cannot be ignored and result is small margin (C= ∞ enforces hard margin)

SVM OPTIMIZATION

- ▶ Constraint can be rewritten as:

$$y_i f(x_i) \geq 1 - \xi_i \quad \forall i$$

- ▶ Together with $\xi_i \geq 0$, is equivalent to:

$$\xi_i = \max\left(0, 1 - y_i f(x_i)\right)$$

- ▶ Hence we can use the following score in unconstrained optimization:

$$\min_{\mathbf{w}} \|\mathbf{w}\|^2 + C \sum_i^N \left[\max\left(0, 1 - y_i f(x_i)\right) \right]$$

NEW OBJECTIVE

$$\min_{\mathbf{w}} \|\mathbf{w}\|^2 + C \sum_i^N \left[\max\left(0, 1 - y_i f(x_i)\right) \right]$$

Hinge Loss

Points are in three categories:

1. $y_i f(x_i) > 1$
Point is outside margin.
No contribution to loss
2. $y_i f(x_i) = 1$
Point is on margin.
No contribution to loss.
As in hard margin case.
3. $y_i f(x_i) < 1$
Point violates margin constraint.
Contributes to loss

